

King County House Price Prediction

In [1]:

```
import numpy as np
import pandas as pd

import seaborn as sns
import matplotlib.pyplot as plt

from scipy import stats
from scipy.stats import boxcox

from sklearn import metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR

from sklearn import ensemble
import xgboost as xgb
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor
import catboost as cb
```

In [2]:

```
data = pd.read_csv('kc_house_data.csv')
df = pd.DataFrame(data)
df.head()
```

Out[2]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
0	7129300520	20141013T000000	221900.0	3	1.00	1180	5650	1.0
1	6414100192	20141209T000000	538000.0	3	2.25	2570	7242	2.0
2	5631500400	20150225T000000	180000.0	2	1.00	770	10000	1.0
3	2487200875	20141209T000000	604000.0	4	3.00	1960	5000	1.0
4	1954400510	20150218T000000	510000.0	3	2.00	1680	8080	1.0

5 rows × 21 columns

In [3]:

```
df.shape
```

Out[3]:

(21613, 21)

In [4]:

```
df.describe()
```

Out[4]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot
count	2.161300e+04	2.161300e+04	21613.000000	21613.000000	21613.000000	2.161300e+04
mean	4.580302e+09	5.400881e+05	3.370842	2.114757	2079.899736	1.510697e+04
std	2.876566e+09	3.671272e+05	0.930062	0.770163	918.440897	4.142051e+04
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21613 non-null  int64
1   date                 21613 non-null  object
2   price               21613 non-null  float64
3   bedrooms            21613 non-null  int64
4   bathrooms           21613 non-null  float64
5   sqft_living         21613 non-null  int64
6   sqft_lot            21613 non-null  int64
7   floors              21613 non-null  float64
8   waterfront          21613 non-null  int64
9   view                21613 non-null  int64
10  condition            21613 non-null  int64
11  grade               21613 non-null  int64
12  sqft_above          21613 non-null  int64
13  sqft_basement       21613 non-null  int64
14  yr_built            21613 non-null  int64
15  yr_renovated        21613 non-null  int64
16  zipcode             21613 non-null  int64
17  lat                 21613 non-null  float64
18  long                21613 non-null  float64
19  sqft_living15       21613 non-null  int64
20  sqft_lot15          21613 non-null  int64
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

NULL Values

In [6]:

```
df.isnull().sum()
```

Out[6]:

id	0
date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
condition	0
grade	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0

dtype: int64


Correlation

In [7]:

```
df.corr()
```

Out[7]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	w
id	1.000000	-0.016762	0.001286	0.005160	-0.012258	-0.132109	0.018525	-
price	-0.016762	1.000000	0.308350	0.525138	0.702035	0.089661	0.256794	
bedrooms	0.001286	0.308350	1.000000	0.515884	0.576671	0.031703	0.175429	-
bathrooms	0.005160	0.525138	0.515884	1.000000	0.754665	0.087740	0.500653	
sqft_living	-0.012258	0.702035	0.576671	0.754665	1.000000	0.172826	0.353949	
sqft_lot	-0.132109	0.089661	0.031703	0.087740	0.172826	1.000000	-0.005201	
floors	0.018525	0.256794	0.175429	0.500653	0.353949	-0.005201	1.000000	
waterfront	-0.002721	0.266369	-0.006582	0.063744	0.103818	0.021604	0.023698	
view	0.011592	0.397293	0.079532	0.187737	0.284611	0.074710	0.029444	
condition	-0.023783	0.036362	0.028472	-0.124982	-0.058753	-0.008958	-0.263768	
grade	0.008130	0.667434	0.356967	0.664983	0.762704	0.113621	0.458183	
sqft_above	-0.010842	0.605567	0.477600	0.685342	0.876597	0.183512	0.523885	
sqft_basement	-0.005151	0.323816	0.303093	0.283770	0.435043	0.015286	-0.245705	
yr_built	0.021380	0.054012	0.154178	0.506019	0.318049	0.053080	0.489319	-
yr_renovated	-0.016907	0.126434	0.018841	0.050739	0.055363	0.007644	0.006338	
zipcode	-0.008224	-0.053203	-0.152668	-0.203866	-0.199430	-0.129574	-0.059121	
lat	-0.001891	0.307003	-0.008931	0.024573	0.052529	-0.085683	0.049614	-
long	0.020799	0.021626	0.129473	0.223042	0.240223	0.229521	0.125419	-
sqft_living15	-0.002901	0.585379	0.391638	0.568634	0.756420	0.144608	0.279885	
sqft_lot15	-0.138798	0.082447	0.029244	0.087175	0.183286	0.718557	-0.011269	



In [8]:

```
df.corr()["price"].sort_values(ascending = True)
```

Out[8]:

zipcode	-0.053203
id	-0.016762
long	0.021626
condition	0.036362
yr_built	0.054012
sqft_lot15	0.082447
sqft_lot	0.089661
yr_renovated	0.126434
floors	0.256794
waterfront	0.266369
lat	0.307003
bedrooms	0.308350
sqft_basement	0.323816
view	0.397293
bathrooms	0.525138
sqft_living15	0.585379
sqft_above	0.605567
grade	0.667434
sqft_living	0.702035
price	1.000000


Name: price, dtype: float64

In [9]:

```
plt.figure(figsize=(18,18))
sns.heatmap(df.corr(),annot=True,cmap="Blues",fmt='.2f',square=True)
```

Out[9]:

<AxesSubplot:>



Feature Engineering


In [10]:

```
df =df.drop(["id"],axis = 1)

df.head()
```

Out[10]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	\
0	20141013T000000	221900.0	3	1.00	1180	5650	1.0	0	
1	20141209T000000	538000.0	3	2.25	2570	7242	2.0	0	
2	20150225T000000	180000.0	2	1.00	770	10000	1.0	0	
3	20141209T000000	604000.0	4	3.00	1960	5000	1.0	0	
4	20150218T000000	510000.0	3	2.00	1680	8080	1.0	0	



In [11]:

```
df['date'] = pd.to_datetime(df['date'])
df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
df['day'] = df['date'].dt.day
df.head()
```

Out[11]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condit
0	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0	0	
1	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0	0	
2	2015-02-25	180000.0	2	1.00	770	10000	1.0	0	0	
3	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0	0	
4	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0	0	

5 rows × 23 columns



In [12]:

```
df = df.drop(['date'],axis = 1)
```


In [13]:

```
df.head()
```

Out[13]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	gr
0	221900.0	3	1.00	1180	5650	1.0	0	0	3	
1	538000.0	3	2.25	2570	7242	2.0	0	0	3	
2	180000.0	2	1.00	770	10000	1.0	0	0	3	
3	604000.0	4	3.00	1960	5000	1.0	0	0	5	
4	510000.0	3	2.00	1680	8080	1.0	0	0	3	

5 rows × 22 columns



In [14]:

```
df['yr_built'].values[1763],df['yr_renovated'].values[1763],df['year'].values[1763]
```

Out[14]:

(2015, 0, 2014)

In [15]:

```
df.loc[df["yr_renovated"] == 0, "yr_renovated"] = df['yr_built']
```

In [16]:

```
df['age'] = (df['year'] - df['yr_renovated'])
```

```
df.head()
```

Out[16]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	gr
0	221900.0	3	1.00	1180	5650	1.0	0	0	3	
1	538000.0	3	2.25	2570	7242	2.0	0	0	3	
2	180000.0	2	1.00	770	10000	1.0	0	0	3	
3	604000.0	4	3.00	1960	5000	1.0	0	0	5	
4	510000.0	3	2.00	1680	8080	1.0	0	0	3	

5 rows × 23 columns



In [17]:

```
df.corr()["price"].sort_values(ascending = False)
```

Out[17]:

```
price          1.000000
sqft_living    0.702035
grade          0.667434
sqft_above     0.605567
sqft_living15  0.585379
bathrooms      0.525138
view           0.397293
sqft_basement  0.323816
bedrooms       0.308350
lat            0.307003
waterfront     0.266369
floors         0.256794
yr_renovated   0.105755
sqft_lot       0.089661
sqft_lot15     0.082447
yr_built       0.054012
condition      0.036362
long           0.021626
year           0.003576
month          -0.010081
day            -0.014670
zipcode        -0.053203
age            -0.105672
Name: price, dtype: float64
```

In [18]:

```
df = df.drop(columns = ['month', 'day', 'year', 'yr_built', 'yr_renovated'], axis = 1)
```

In [19]:

```
index=[]
x=0
for i in df['age'].values:
    if i<0:
        print(x,i)
        index.append(x)
    x+=1
```

```
1763 -1
2295 -1
2687 -1
7097 -1
7526 -1
8039 -1
11599 -1
14489 -1
14859 -1
15687 -1
17098 -1
18575 -1
19805 -1
20770 -1
20852 -1
20963 -1
21262 -1
21372 -1
```

In [20]:

```
# The age is in negative, because year_sold < year_built which is not possible. So dropping
```

In [21]:

```
df = df.drop(index,axis = 0)
df = df.drop(columns = ['zipcode'],axis = 1)
```

In [22]:

```
df.head()
```

Out[22]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	gr
0	221900.0	3	1.00	1180	5650	1.0	0	0	3	
1	538000.0	3	2.25	2570	7242	2.0	0	0	3	
2	180000.0	2	1.00	770	10000	1.0	0	0	3	
3	604000.0	4	3.00	1960	5000	1.0	0	0	5	
4	510000.0	3	2.00	1680	8080	1.0	0	0	3	



Cheking Skewness

In [23]:

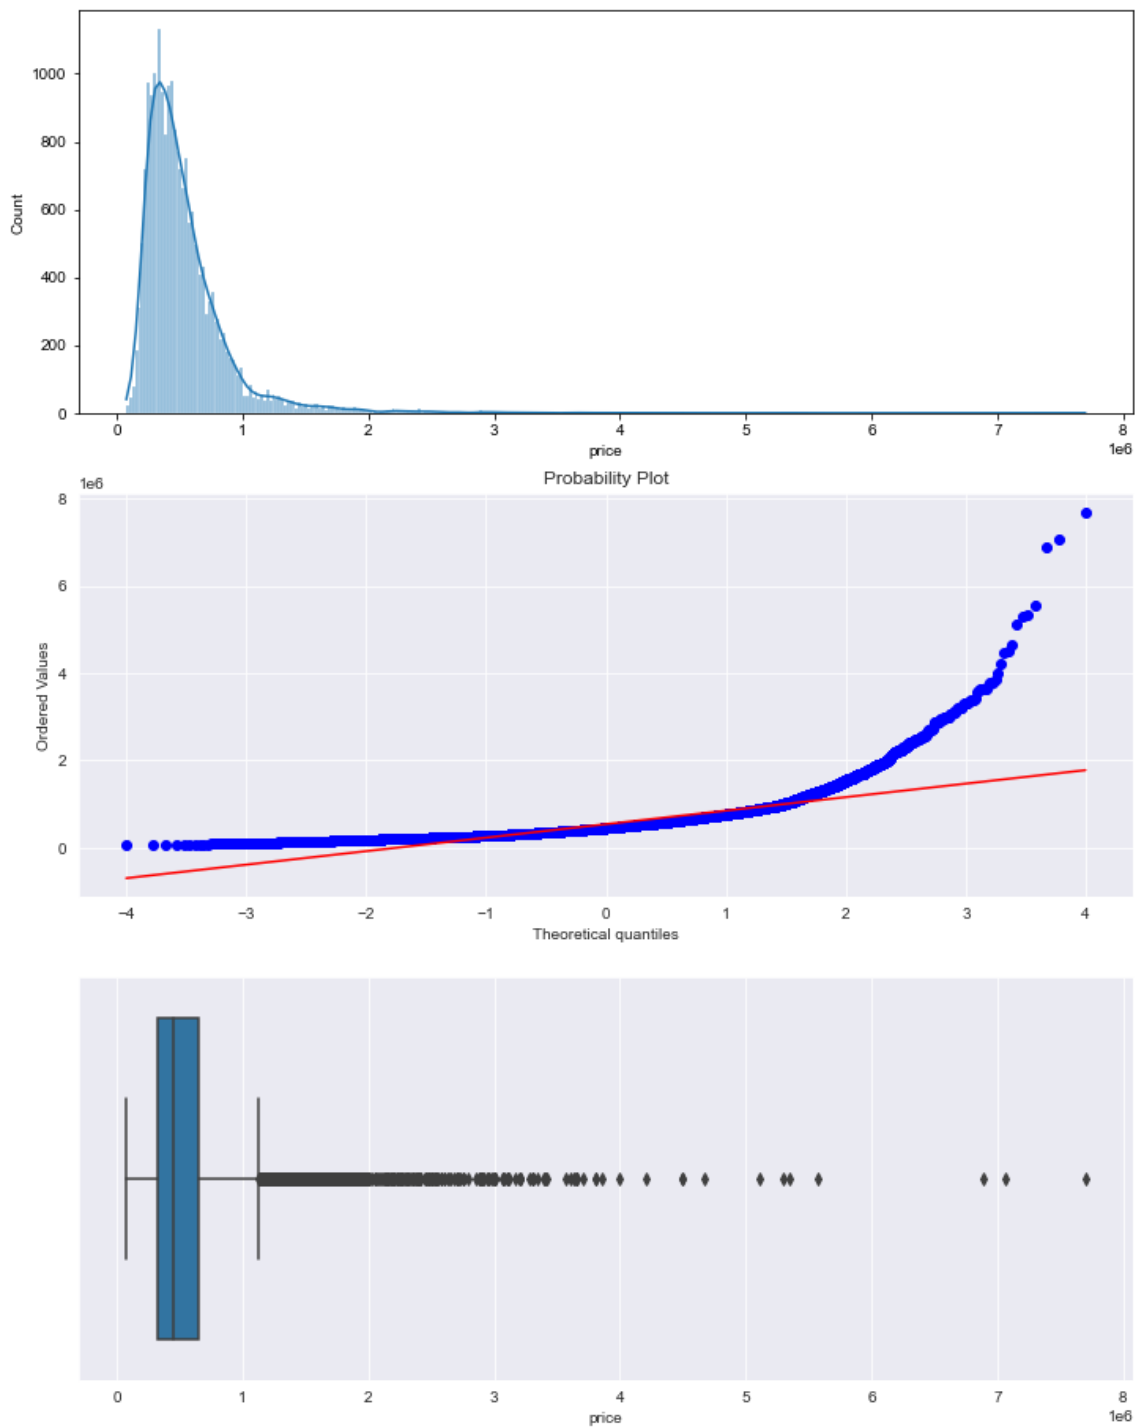
```
#Checking Skewness
fig=plt.figure(figsize=(12,16))
fig.tight_layout()
ax1 = fig.add_subplot(3, 1, 1)
sns.set_style("darkgrid")
sns.histplot(df.loc[:, 'price'],kde=True,ax=ax1)

ax2 = fig.add_subplot(3, 1, 2)
stats.probplot(df.loc[:, 'price'],plot=ax2)

ax3 = fig.add_subplot(3, 1, 3)
sns.boxplot(x=df.loc[:, 'price'],ax=ax3)
```

Out[23]:

<AxesSubplot:xlabel='price'>



In [24]:

```
df['price']=boxcox(df['price'])[0]
```

In [25]:

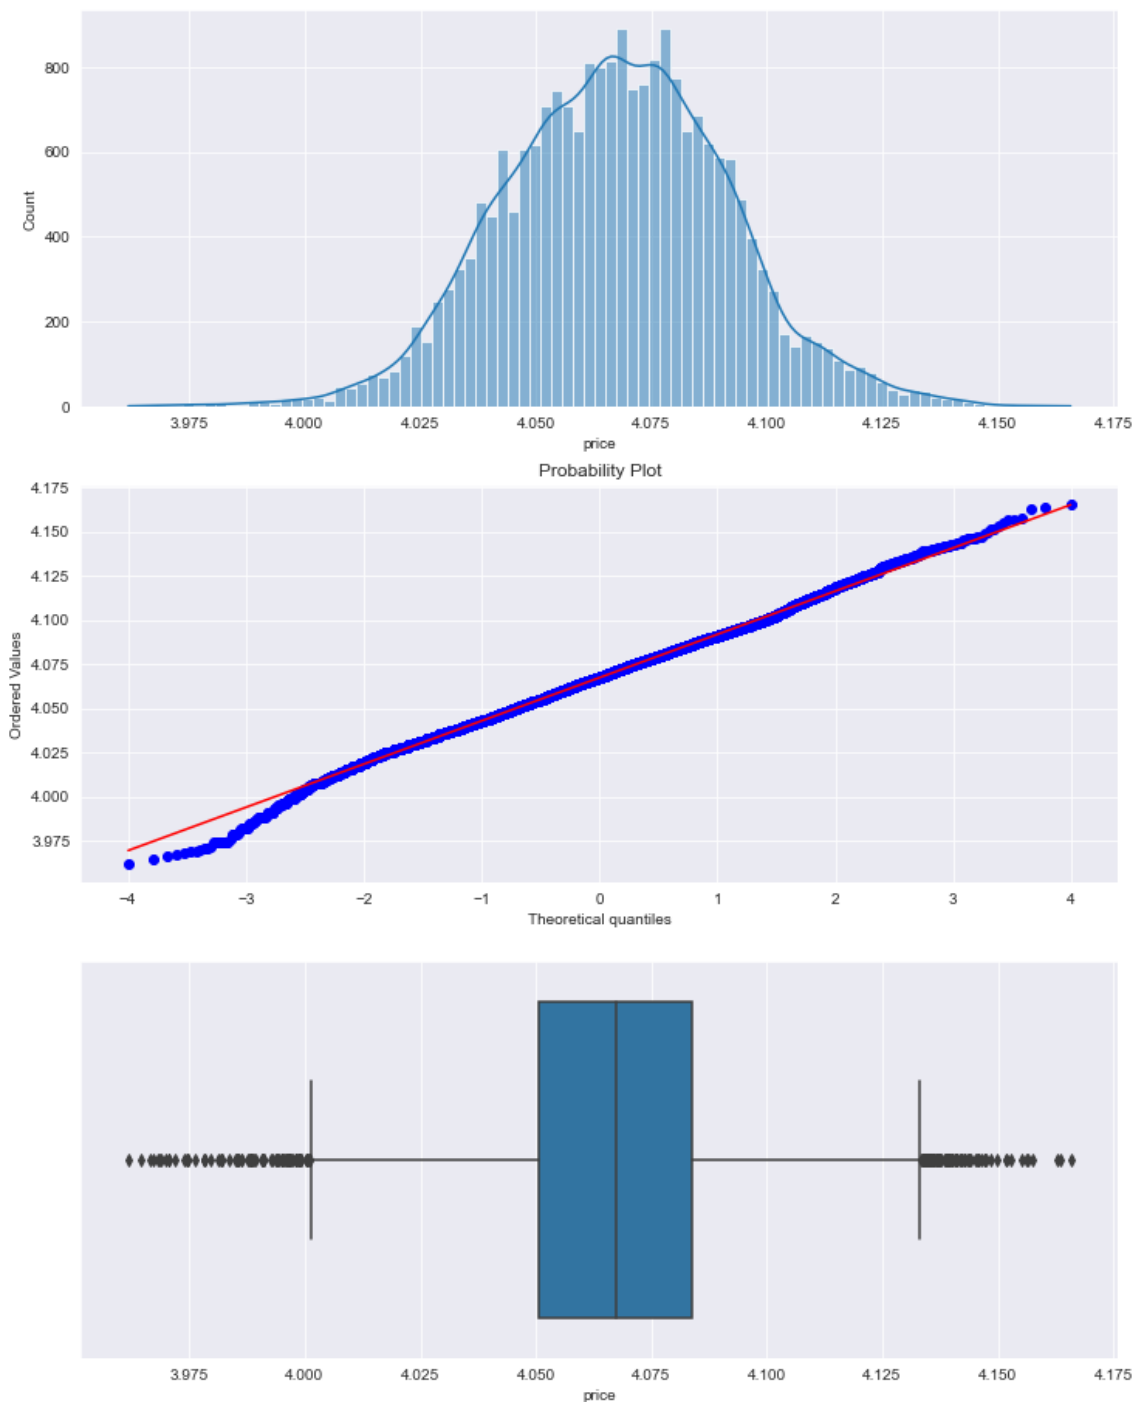
```
#Checking Skewness
fig=plt.figure(figsize=(12,16))
fig.tight_layout()
ax1 = fig.add_subplot(3, 1, 1)
sns.set_style("darkgrid")
sns.histplot(df.loc[:, 'price'],kde=True,ax=ax1)

ax2 = fig.add_subplot(3, 1, 2)
stats.probplot(df.loc[:, 'price'],plot=ax2)

ax3 = fig.add_subplot(3, 1, 3)
sns.boxplot(x=df.loc[:, 'price'],ax=ax3)
```

Out[25]:

<AxesSubplot:xlabel='price'>



In [26]:

```
df.shape
```

Out[26]:

```
(21595, 17)
```

In [27]:

```
i=0
while i<df.shape[1]:
    x=df.columns[i]
    if df[x].dtypes!=object and abs(df[x]).skew() >0.5 and i!=0:
        df[x]=stats.boxcox(df[x]+1)[0]
    i+=1
```

In [28]:

```
df
```

Out[28]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	cond
0	4.030919	1.845296	0.790440	7.814065	4.195162	0.433862	0.0	0.0	0.80
1	4.075683	1.845296	1.478450	8.770486	4.240651	0.543474	0.0	0.0	0.80
2	4.018914	1.375287	0.790440	7.298470	4.296609	0.433862	0.0	0.0	0.80
3	4.080876	2.248442	1.812704	8.435165	4.171951	0.433862	0.0	0.0	0.91
4	4.073237	1.845296	1.356321	8.245539	4.260032	0.433862	0.0	0.0	0.80
...
21608	4.056522	1.845296	1.594803	8.130891	3.841087	0.589749	0.0	0.0	0.80
21609	4.061723	2.248442	1.594803	8.638189	4.200485	0.543474	0.0	0.0	0.80
21610	4.061978	1.375287	0.621954	7.637361	3.885643	0.543474	0.0	0.0	0.80
21611	4.061723	1.845296	1.594803	8.185693	4.019296	0.543474	0.0	0.0	0.80
21612	4.051349	1.375287	0.621954	7.637361	3.828261	0.543474	0.0	0.0	0.80

21595 rows × 10 columns

Scaling

In [29]:

```
df_scaled=df
from sklearn.preprocessing import StandardScaler
#scaled=RobustScaler()
scaled = StandardScaler()

df_scaled = pd.DataFrame(scaled.fit_transform(df),columns=df.columns)

df_scaled.head()
```

Out[29]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	co
0	-1.487163	-0.352330	-1.571621	-1.121672	-0.332344	-0.975183	-0.087209	-0.329935	-0.
1	0.341151	-0.352330	0.251450	0.706994	-0.032751	1.030360	-0.087209	-0.329935	-0.
2	-1.977451	-1.606097	-1.571621	-2.107485	0.335789	-0.975183	-0.087209	-0.329935	-0.
3	0.553232	0.723079	1.137145	0.065865	-0.485213	-0.975183	-0.087209	-0.329935	2.
4	0.241236	-0.352330	-0.072163	-0.296698	0.094890	-0.975183	-0.087209	-0.329935	-0.

Dummies

In [30]:

```
features = pd.get_dummies(df_scaled).reset_index(drop=True)
features.shape
```

Out[30]:

(21595, 17)

In [31]:

```
X = df_scaled.drop(['price'],axis = 1)
X.head()
```

Out[31]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	
0	-0.352330	-1.571621	-1.121672	-0.332344	-0.975183	-0.087209	-0.329935	-0.598191	-0.
1	-0.352330	0.251450	0.706994	-0.032751	1.030360	-0.087209	-0.329935	-0.598191	-0.
2	-1.606097	-1.571621	-2.107485	0.335789	-0.975183	-0.087209	-0.329935	-0.598191	-1.
3	0.723079	1.137145	0.065865	-0.485213	-0.975183	-0.087209	-0.329935	2.080336	-0.
4	-0.352330	-0.072163	-0.296698	0.094890	-0.975183	-0.087209	-0.329935	-0.598191	0.

In [32]:

```
y = df_scaled['price']
y.head()
```

Out[32]:

```
0    -1.487163
1     0.341151
2    -1.977451
3     0.553232
4     0.241236
Name: price, dtype: float64
```

In [33]:

```
rms=[]
mse=[]
mae=[]
r2=[]

def evaluation(model, x_test, y_test):
    pred = model.predict(x_test)

    rms.append(np.sqrt(mean_squared_error(y_test, pred)))

    mse.append(mean_squared_error(y_test, pred))

    mae.append(mean_absolute_error(y_test, pred))

    r2.append(r2_score(y_test,pred))

    print("Result\n", "-----")
    print("\nRoot Mean Squared Error: ",np.sqrt(mean_squared_error(y_test, pred)))
    print("\nMean Squared Error: \t",mean_squared_error(y_test, pred))
    print("\nMean Absolute Error: \t",mean_absolute_error(y_test, pred))
    print("\nR2 Score: \t\t",r2_score(y_test,pred))
```

Splitting data for training and testing

In [34]:

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

Linear Regression

In [35]:

```
lg = LinearRegression()

model=lg.fit(X_train,y_train)

y_pred_lg = lg.predict(X_test)

y_pred_lg
```

Out[35]:

```
array([-1.34381624, -0.15454372, -0.02232806, ..., -0.0058335 ,
       -0.83150375, -0.2049119 ])
```

In [36]:

```
evaluation(model, X_test, y_test)
```

Result

```
-----  
Root Mean Squared Error:  0.5107123586919792  
Mean Squared Error:      0.26082711332072483  
Mean Absolute Error:     0.391958205061279  
R2 Score:                 0.7393462089171894
```

KNN

In [37]:

```
kr = KNeighborsRegressor(n_neighbors=10)  
model=kr.fit(X_train,y_train)
```

In [38]:

```
evaluation(model, X_test, y_test)
```

Result

```
-----  
Root Mean Squared Error:  0.42202253003026213  
Mean Squared Error:      0.1781030158531435  
Mean Absolute Error:     0.30242172317552296  
R2 Score:                 0.8220153353906897
```

SVR

In [39]:

```
sv = SVR(kernel = 'rbf')  
model=sv.fit(X_train, y_train)
```

In [40]:

```
evaluation(model, X_test, y_test)
```

Result

```
-----  
Root Mean Squared Error:  0.38780732856598055  
  
Mean Squared Error:      0.1503945240894824  
  
Mean Absolute Error:     0.2730300617046993  
  
R2 Score:                 0.8497054145831248
```

XGB

In [41]:

```
xgb_regress = xgb.XGBRegressor(n_estimators = 2000, learning_rate = 0.1)  
l=xgb_regress.fit(X_train, y_train)  
  
y_pred_xg = l.predict(X_test)  
accuracy = round(r2_score(y_test,y_pred_xg),3)  
accuracy
```

Out[41]:

0.879

In [42]:

```
evaluation(l, X_test, y_test)
```

Result

```
-----  
Root Mean Squared Error:  0.3481417810946497  
  
Mean Squared Error:      0.12120269974375499  
  
Mean Absolute Error:     0.23691109537962213  
  
R2 Score:                 0.878877840668219
```

Gradient Boosting

In [43]:

```
params = {  
    "n_estimators": 1000,  
    "max_depth": 4,  
    "min_samples_split": 5,  
    "learning_rate": 0.1  
}  
  
reg = ensemble.GradientBoostingRegressor(**params)  
l=reg.fit(X_train, y_train)
```

In [44]:

```
evaluation(l, X_test, y_test)
```

Result

```
-----  
  
Root Mean Squared Error:  0.3397205518739977  
  
Mean Squared Error:      0.11541005336557356  
  
Mean Absolute Error:     0.23384886750326775  
  
R2 Score:                 0.8846666377746705
```

ADABOOST

In [45]:

```
regr = ensemble.AdaBoostRegressor(random_state=0, n_estimators=1000)  
l=regr.fit(X_train, y_train)
```

In [46]:

```
evaluation(l, X_test, y_test)
```

Result

```
-----  
  
Root Mean Squared Error:  0.4985404861954686  
  
Mean Squared Error:      0.24854261637601419  
  
Mean Absolute Error:     0.38537733668965807  
  
R2 Score:                 0.7516225426902305
```

Random Forest

In [47]:

```
regressor = RandomForestRegressor(n_estimators = 100, random_state = 0)

# fit the regressor with x and y data
l=regressor.fit(X_train, y_train)
```

In [48]:

```
evaluation(l, X_test, y_test)
```

Result

Root Mean Squared Error: 0.36239858625509846

Mean Squared Error: 0.13133273531969406

Mean Absolute Error: 0.24957256832632102

R2 Score: 0.8687545366027192

CAT Boost

In [49]:

```
model = cb.CatBoostRegressor(iterations=1000,
                             learning_rate=0.1,
                             depth=8)

# Fit model
cat=model.fit(X_train, y_train)
```

0:	learn: 0.9314973	total: 156ms	remaining: 2m 35s
1:	learn: 0.8703926	total: 164ms	remaining: 1m 22s
2:	learn: 0.8132449	total: 173ms	remaining: 57.3s
3:	learn: 0.7617152	total: 181ms	remaining: 45.1s
4:	learn: 0.7153295	total: 189ms	remaining: 37.7s
5:	learn: 0.6777702	total: 197ms	remaining: 32.7s
6:	learn: 0.6426906	total: 205ms	remaining: 29.1s
7:	learn: 0.6101366	total: 213ms	remaining: 26.5s
8:	learn: 0.5829402	total: 221ms	remaining: 24.3s
9:	learn: 0.5584274	total: 229ms	remaining: 22.7s
10:	learn: 0.5363441	total: 237ms	remaining: 21.3s
11:	learn: 0.5171809	total: 246ms	remaining: 20.3s
12:	learn: 0.4997535	total: 255ms	remaining: 19.4s
13:	learn: 0.4853623	total: 264ms	remaining: 18.6s
14:	learn: 0.4692499	total: 272ms	remaining: 17.9s
15:	learn: 0.4555180	total: 280ms	remaining: 17.2s
16:	learn: 0.4451452	total: 288ms	remaining: 16.6s
17:	learn: 0.4350465	total: 296ms	remaining: 16.2s
18:	learn: 0.4267540	total: 304ms	remaining: 15.7s
19:	learn: 0.4183350	total: 312ms	remaining: 15.2s

In [50]:

```
evaluation(cat, X_test, y_test)
```

Result

Root Mean Squared Error: 0.33141486335313103

Mean Squared Error: 0.1098358116513745

Mean Absolute Error: 0.22662380597553808

R2 Score: 0.8902371753492335

Results

In [51]:

```
res=pd.DataFrame()  
res['Model']=["LR","KNN","SVM","XGB","GRA","ADA","RF","CAT"]  
res['RMSE']=rms  
res['MSE']=mse  
res['MAE']=mae  
res['R2']=r2
```

In [52]:


```
res
```

Out[52]:

	Model	RMSE	MSE	MAE	R2
0	LR	0.510712	0.260827	0.391958	0.739346
1	KNN	0.422023	0.178103	0.302422	0.822015
2	SVM	0.387807	0.150395	0.273030	0.849705
3	XGB	0.348142	0.121203	0.236911	0.878878
4	GRA	0.339721	0.115410	0.233849	0.884667
5	ADA	0.498540	0.248543	0.385377	0.751623
6	RF	0.362399	0.131333	0.249573	0.868755
7	CAT	0.331415	0.109836	0.226624	0.890237

In [53]:

```
sns.set(style="whitegrid")  
ax=sns.barplot(y='Model',x='R2',data=res)
```



In []: